

PRACTICAL JOURNAL IN
APPLIED ARTIFICIAL INTELLIGENCE
MACHINE LEARNING

SUBMITTED BY:

SHUBHAM LAXMAN PATTADE

ROLL NO: 2414568

IN PARTIAL FULLFILMENT FOR THE DEGREE OF
MASTER OF SCIENCE IN INFORMATION TECHNOLOGY PART – II
SEMESTER III

ACADEMIC YEAR
2024-2025



PARLE TILAK VIDYALAYA ASSOCIATION'S
MULUND COLLEGE OF COMMERCE(AUTONOMOUS)
(AFFILIATED TO UNIVERSITY OF MUMBAI)
NAAC RE-ACCREDITED A GRADE – III CYCLE
MULUND WEST, MUMBAI 400080
MAHARASHTRA, INDIA
2024-25



Parle Tilak Vidyalaya Association's
MULUND COLLEGE OF COMMERCE
(AUTONOMOUS)

(Affiliated to University of Mumbai)
MULUND WEST, MUMBAI 400080
MAHARASHTRA, INDIA

DEPARTMENT OF INFORMATION TECHNOLOGY

CERTIFICATE

This is to certify that **Shubham L Pattade** of **M.Sc. I.T. Part II** Roll No **2414568** has successfully completed the practical work in **Applied Artificial Intelligence** in partial fulfilment of the requirements for the Semester III of **M.Sc. I.T. Part II** during the academic year **2024-25**.

Teacher In-charge and Coordinator

Examiner

Date:

College Seal



Parle Tilak Vidyalaya Association's
MULUND COLLEGE OF COMMERCE
(AUTONOMOUS)

(Affiliated to University of Mumbai)
MULUND WEST, MUMBAI 400080
MAHARASHTRA, INDIA

DEPARTMENT OF INFORMATION TECHNOLOGY

CERTIFICATE

This is to certify that **Shubham L Pattade** of **M.Sc. I.T. Part II** Roll No **2414568** has successfully completed the practical work in **Machine Learning** in partial fulfilment of the requirements for the Semester III of **M.Sc. I.T. Part II** during the academic year **2024-25**.

Teacher In-charge and Coordinator

Examiner

Date:

College Seal

Table of Contents

Practical 1: Design a bot using AIML3

Practical 2: Design an Expert system using AIML6

Practical 3: Implement Bayes Theorem using Python10

Practical 4: Implement Conditional Probability and joint probability using Python.11

Practical 5: Write a program to implement Rule based system. (Prolog)12

Practical 6: Design a Fuzzy based application using Python / R17

Practical 7A: Write an application to stimulate supervised learning model.21

7B: Write an application to stimulate unsupervised learning model.23

Practical 8: Write an application to implement clustering algorithm.25

Practical 9: Write an application to implement support vector machine algorithm.27

Practical 10: Simulate artificial neural network model with both feedforward and backpropagation approach.30

Practical 1: Design a bot using AIML

Std-startup.xml

```
<aiml version="1.0.1" encoding="UTF-8">
  <!-- std-startup.xml -->
  <!-- Category is an atomic AIML unit -->
  <category>
    <!-- Pattern to match in user input -->
    <!-- If user enters "LOAD AIML B" -->
    <pattern>LOAD AIML B</pattern>
    <!-- Template is the response to the pattern -->
    <!-- This learn an aiml file -->
    <template>
      <learn>basic_chat.aiml</learn>
      <!-- You can add more aiml files here -->
      <!--<learn>more_aiml.aiml</learn>-->
    </template>
  </category>
</aiml>
```

Basic_chat.aiml

```
<aiml version="1.0.1" encoding="UTF-8">
  <!-- basic_chat.aiml -->
  <category>
    <pattern>HELLO</pattern>
    <template>
      Well, hello!
    </template>
  </category>
  <category>
    <pattern>WHAT ARE YOU</pattern>
    <template> I'm a bot, silly! </template>
  </category>
  <category>
```

```

<pattern>MY NAME IS *</pattern>
    <template>
        <set name="username">
            <star />
        </set> is the nice name. </template>
</category>
<category>
    <pattern>I LIKE *</pattern>
    <template>
        <set name="liking">
            <star />
        </set> is also my favourite. </template>
</category>
<category>
    <pattern>MY DOG NAME IS *</pattern>
    <template> THAT IS INTERESTING THAT YOU HAVE A DOG NAMED <set name="dog">
        <star />
    </set> . </template>
</category>
<category>
    <pattern>BYE</pattern>
    <template> Bye!!! <get name="username" /> Thanks for talking with me. </template>
</category>
</aiml>

```

Chatbot.py

```

import aiml # Create the kernel and learn AIML files
kernel = aiml.Kernel()
kernel.learn("std-startup.xml")
kernel.respond("load aiml b") # Press CTRL-C to break this loop
while True:
    message = input("Enter your message to the bot: ")
    if message == "quit":
        break

```

else:

```
    bot_response = kernel.respond(message)
    print(bot_response)
```

Output:

```
----- RESTART: E:\ABC\Jupyter\AI_1\AI_1\day/ChatBot.py -----
Loading std-startup.xml...done (0.23 seconds)
Loading basic_chat.aiml...done (0.00 seconds)
Enter your message to the bot: hello
Well, hello!
Enter your message to the bot: what are you
I'm a bot, silly!
Enter your message to the bot: quit
```

Practical 2: Design an Expert system using AIML

Std_startup.xml

```
<aiml version="1.0.1" encoding="UTF-8">
  <!-- std-startup.xml -->
  <!-- Category is an atomic AIML unit -->
  <category>
    <!-- Pattern to match in user input -->
    <!-- If user enters "LOAD AIML B" -->
    <pattern>LOAD AIML B</pattern>
    <!-- Template is the response to the pattern -->
    <!-- This learn an aiml file -->
    <template>
      <learn>basic_chat.aiml</learn>
      <!-- You can add more aiml files here -->
      <!--<learn>more_aiml.aiml</learn>-->
    </template>
  </category>
</aiml>
```

Basic_chat.aiml

```
<aiml version="1.0.1" encoding="UTF-8">
  <!-- basic_chat.aiml -->
  <category>
    <pattern>HELLO</pattern>
    <template>
      WHAT WOULD YOU LIKE TO DISCUSS? : HEALTH, MOVIES
    </template>
  </category>
  <category>
    <pattern>MOVIES</pattern>
    <template>
      YES <set name = "topic">MOVIES</set>
    </template>
```



```

</category>
  <category>
    <pattern>HEALTH</pattern>
    <template> YES <set name = "topic">HEALTH</set> </template>
  </category>
  <topic name ="MOVIES">
    <category>
      <pattern>*</pattern>
      <template>
        DO YOU LIKE COMEDY MOVIES?
      </template>
    </category>
    <category> <pattern>YES</pattern>
    <template>
      I TOO LIKE COMEDY MOVIES
    </template>
  </category>
  <category>
    <pattern>NO</pattern>
    <template>
      BUT I LIKE COMEDY MOVIES
    </template>
  </category>
</topic>
<topic name ="HEALTH">
  <category>
    <pattern>*</pattern>
    <template>
      DO YOU HAVE FEVER?
    </template>
  </category>
  <category>
    <pattern>YES</pattern>
    <template>

```

PLEASE TAKE MEDICINES AND PROPER REST

</template></category>

<category>

<pattern>NO</pattern>

<template>

GO OUT FOR A WALK AND LISTEN MUSIC

</template>

</category>

</topic>

<category>

<pattern>NICE TALKING TO YOU</pattern>

<template>

SAME HERE....!!

</template>

</category>

</aiml>

Chatbot.py

```
import aiml

# Create the kernel and learn AIML files
kernel = aiml.Kernel()
kernel.learn("std-startup.xml")
kernel.respond("LOAD AIML B")
# Press CTRL-C to break this loop
while True:
    message = input("Enter your message to the bot: ")
    if message == "quit":
        break
    else:
        bot_response = kernel.respond(message)
        print(bot_response)
```

Output:

```
----- RESIARI. E./MSC/SEMS/AA1 F1aC/F1a1ay/F2.py -----
Loading std-startup.xml...done (0.03 seconds)
Loading basic_chat.aiml...done (0.00 seconds)
Enter your message to the bot: hello
WHAT WOULD YOU LIKE TO DISCUSS? : HEALTH, MOVIES
Enter your message to the bot: health
YES HEALTH
Enter your message to the bot: i am feeling tired
DO YOU HAVE FEVER?
Enter your message to the bot: no
GO OUT FOR A WALK AND LISTEN MUSIC
Enter your message to the bot: Movies
YES MOVIES
Enter your message to the bot: I love movies
DO YOU LIKE COMEDY MOVIES?
Enter your message to the bot: yes
I TOO LIKE COMEDY MOVIES
Enter your message to the bot: Thanks
DO YOU LIKE COMEDY MOVIES?
Enter your message to the bot: Quit
DO YOU LIKE COMEDY MOVIES?
Enter your message to the bot: yes
I TOO LIKE COMEDY MOVIES
Enter your message to the bot: Yes
I TOO LIKE COMEDY MOVIES
Enter your message to the bot: quit
|
```

Practical 3: Implement Bayes Theorem using Python

```
# calculate the probability of cancer patient and diagnostic test
# calculate P(A|B) given P(A), P(B|A), P(B|not A)
def bayes_theorem(p_a, p_b_given_a, p_b_given_not_a):
    # calculate P(not A)
    not_a = 1 - p_a
    # calculate P(B)
    p_b = p_b_given_a * p_a + p_b_given_not_a * not_a
    # calculate P(A|B)
    p_a_given_b = (p_b_given_a * p_a) / p_b
    return p_a_given_b

# P(A)
p_a = 0.0002
# P(B|A)
p_b_given_a = 0.85
# P(B|not A)
p_b_given_not_a = 0.05
# calculate P(A|B)
result = bayes_theorem(p_a, p_b_given_a, p_b_given_not_a)
# summarize
print('P(A|B) = %.3f%%' % (result * 100))
```

Output:

```
===== RESTART: E:\MSC\Sem3\AAI Prac\Pranay\P3.py =====
P(A|B) = 0.339%
```

Practical 4: Implement Conditional Probability and joint probability using Python.

```
import enum, random

class Kid(enum.Enum):
    BOY = 0
    GIRL = 1

def random_kid() -> Kid:
    return random.choice([Kid.BOY, Kid.GIRL])

both_girls = 0
older_girl = 0
either_girl = 0
random.seed(0)

for _ in range(10000):
    younger = random_kid()
    older = random_kid()
    if older == Kid.GIRL:
        older_girl += 1
    if older == Kid.GIRL and younger == Kid.GIRL:
        both_girls += 1
    if older == Kid.GIRL or younger == Kid.GIRL:
        either_girl += 1

print("older girl: ", older_girl)
print("both girl: ", both_girls)
print("either girl: ", either_girl)
print("P(both | older):", both_girls / older_girl)
print("P(both | either):", both_girls / either_girl)
```

Output:

```
===== RESTART: E:\MSC\SEMS\AA1 PTA\Pranay\F4.py =====
older girl:  4937
both girl:   2472
either girl:  7464
P(both | older): 0.5007089325501317
P(both | either): 0.3311897106109325
```

Practical 5: Write a program to implement Rule based system. (Prolog)

```
go:-
hypothesis(Disease),
write('I believe that the patient have '),
write(Disease),
nl,
write('TAKE CARE '),
undo.

/*Hypothesis that should be tested*/
hypothesis(cold) :- cold, !.
hypothesis(flu) :- flu, !.
hypothesis(typhoid) :- typhoid, !.
hypothesis(measles) :- measles, !.
hypothesis(malaria) :- malaria, !.
hypothesis(unknown). /* no diagnosis*/

/*Hypothesis Identification Rules*/
cold :-
verify(headache),
verify(runny_nose),
verify(sneezing),
verify(sore_throat),
write('Advices and Sugestions:'),
nl,
write('1: Tylenol/tab'),
nl,
write('2: panadol/tab'),
nl,
write('3: Nasal spray'),
nl,
write('Please wear warm cloths Because'),
nl.
```

flu :-

```
verify(fever),
verify(headache),
verify(chills),
verify(body_ache),
write('Advices and Sugestions:'),
nl,
write('1: Tamiflu/tab'),
nl,
write('2: panadol/tab'),
nl,
write('3: Zanamivir/tab'),
nl,
write('Please take a warm bath and do salt gargling Because'),
nl.
```

typhoid :-

```
verify(headache),
verify(abdominal_pain),
verify(poor_appetite),
verify(fever),
write('Advices and Sugestions:'),
nl,
write('1: Chloramphenicol/tab'),
nl,
write('2: Amoxicillin/tab'),
nl,
write('3: Ciprofloxacin/tab'),
nl,
write('4: Azithromycin/tab'),
nl,
write('Please do complete bed rest and take soft Diet Because'),
nl.
```

measles :-

```
verify(fever),
```

verify(runny_nose),

verify(rash),

verify(conjunctivitis),

write('Advices and Sugestions:'),

nl,

write('1: Tylenol/tab'),

nl,

write('2: Aleve/tab'),

nl,

write('3: Advil/tab'),

nl,

write('4: Vitamin A'),

nl,

write('Please Get rest and use more liquid Because'),

nl.

malaria :-

verify(fever),

verify(sweating),

verify(headache),

verify(nausea),

verify(vomiting),

verify(diarrhea),

write('Advices and Sugestions:'),

nl,

write('1: Aralen/tab'),

nl,

write('2: Qualaquin/tab'),

nl,

write('3: Plaquenil/tab'),

nl,

write('4: Mefloquine'),

nl,

write('Please do not sleep in open air and cover your full skin Because'),

nl.


```
/* how to ask questions */
```

```
ask(Question) :-
```

```
write('Does the patient have following symptom:'),
```

```
write(Question),
```

```
write('? '),
```

```
read(Response),
```

```
nl,
```

```
( (Response == yes ; Response == y)
```

```
->
```

```
assert(yes(Question)) ;
```

```
assert(no(Question)), fail).
```

```
:- dynamic yes/1,no/1.
```

```
/*How to verify something */
```

```
verify(S) :-
```

```
(yes(S)
```

```
->
```

```
true ;
```

```
(no(S)
```

```
->
```

```
fail ;
```

```
ask(S))).
```

```
/* undo all yes/no assertions*/
```

```
undo :- retract(yes(_)),fail.
```

```
undo :- retract(no(_)),fail.
```

```
undo.
```

Output:

```
?- go.  
Does the patient have following symptom:headache? y.  
  
Does the patient have following symptom:runny_nose? |: y.  
  
Does the patient have following symptom:sneezing? |: y.  
  
Does the patient have following symptom:sore_throat? |: y.  
  
Advices and Sugestions:  
1: Tylenol/tab  
2: panadol/tab  
3: Nasal spray  
Please wear warm cloths Because  
I believe that the patient have cold  
TAKE CARE  
true.  
  
?- ■
```

Practical 6: Design a Fuzzy based application using Python / R

```
import numpy as np
import skfuzzy as fuzz
import matplotlib.pyplot as plt
from skfuzzy import control as ctrl
from mpl_toolkits.mplot3d import Axes3D # Required for 3D plotting
# New Antecedent/Consequent objects hold universe variables and membership
# functions
quality = ctrl.Antecedent(np.arange(0, 10, 0.1), 'quality')
service = ctrl.Antecedent(np.arange(0, 10, 0.1), 'service')
tip = ctrl.Consequent(np.arange(0, 25, 0.1), 'tip')

quality['poor'] = fuzz.zmf(quality.universe, 0,5)
quality['average'] = fuzz.gaussmf(quality.universe,5,1)
quality['good'] = fuzz.smf(quality.universe,5,10)

service['poor'] = fuzz.zmf(service.universe, 0,5)
service['average'] = fuzz.gaussmf(service.universe,5,1)
service['good'] = fuzz.smf(service.universe,5,10)

tip['low'] = fuzz.trimf(tip.universe, [0, 0, 13])
tip['medium'] = fuzz.trimf(tip.universe, [0, 13, 25])
tip['high'] = fuzz.trimf(tip.universe, [13, 25, 25])

quality['average'].view()
plt.title('Quality')



service['poor'].view()
```


```
plt.title('Service')

tip['medium'].view()
plt.title('Tip Medium')

rule1 = ctrl.Rule(quality['poor'] | service['poor'], tip['low'])
rule2 = ctrl.Rule(service['average'], tip['medium'])
rule3 = ctrl.Rule(service['good'] | quality['good'], tip['high'])
rule1.view()
plt.title('Rule 1')
rule2.view()
plt.title('Rule 2')
rule3.view()
plt.title('Rule 3')
tipping_ctrl = ctrl.ControlSystem([rule1, rule2, rule3])
tipping = ctrl.ControlSystemSimulation(tipping_ctrl)
tipping.input['quality'] = 6.5
tipping.input['service'] = 9.8
tipping.compute()
plt.title('Result')
plt.show(block=True)
print(tipping.output['tip'])
tip.view(sim=tipping)
```

Output:





Practical 7A: Write an application to stimulate supervised learning model.

```
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix
from sklearn import datasets

iris=datasets.load_iris()

x = iris.data
y = iris.target

print ('sepal-length', 'sepal-width', 'petal-length', 'petal-width')
print(x)
print('class: 0-Iris-Setosa, 1- Iris-Versicolour, 2- Iris-Virginica')
print(y)

x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)
#To Training the model and Nearest neighbors K=5
classifier = KNeighborsClassifier(n_neighbors=5)
classifier.fit(x_train, y_train)

#To make predictions on our test data
y_pred=classifier.predict(x_test)
print('Confusion Matrix')
print(confusion_matrix(y_test,y_pred))
print('Accuracy Metrics')
print(classification_report(y_test,y_pred))
```

[illegible]

7B: Write an application to stimulate unsupervised learning model.

```
from scipy.cluster.hierarchy import linkage, dendrogram
import matplotlib.pyplot as plt
import pandas as pd

# Reading the DataFrame
seeds_df = pd.read_csv("seeds-less-rows.csv")
# Remove the grain species from the DataFrame, save for later
varieties = list(seeds_df.pop('grain_variety'))
# Extract the measurements as a NumPy array
samples = seeds_df.values

"""
Perform hierarchical clustering on samples using the
linkage() function with the method='complete' keyword argument.
Assign the result to mergings.
"""


mergings = linkage(samples, method='complete')

"""
Plot a dendrogram using the dendrogram() function on mergings,
specifying the keyword arguments labels=varieties, leaf_rotation=90,
and leaf_font_size=6.
"""

dendrogram(mergings,
            labels=varieties,
            leaf_rotation=90,
            leaf_font_size=6,
            )

plt.show()
```

Output:



Practical 8: Write an application to implement clustering algorithm.

```
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.cluster import KMeans
import sklearn.metrics as sm
import pandas as pd
import numpy as np


iris = datasets.load_iris()
X = pd.DataFrame(iris.data)
X.columns = ['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width']
y = pd.DataFrame(iris.target)
y.columns = ['Targets']
model = KMeans(n_clusters=3)
model.fit(X)

plt.figure(figsize=(14,7))
colormap = np.array(['red', 'lime', 'black'])
# Plot the Original Classifications
plt.subplot(1, 2, 1)
plt.scatter(X.Petal_Length, X.Petal_Width,
c=colormap[y.Targets], s=40)
plt.title('Real Classification')
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')
# Plot the Models Classifications
plt.subplot(1, 2, 2)
plt.scatter(X.Petal_Length, X.Petal_Width,
c=colormap[model.labels_], s=40)
plt.title('K Mean Classification')
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')
plt.show()
```

```
print('The accuracy score of K-Mean: ',sm.accuracy_score(y, model.labels_))
```

```
print('The Confusion matrix of K-Mean: ',sm.confusion_matrix(y, model.labels_))
```

Output:



Practical 9: Write an application to implement support vector machine algorithm.

```
#Import scikit-learn dataset library
from sklearn import datasets

#Import svm model
from sklearn import svm

# Import train_test_split function
from sklearn.model_selection import train_test_split

#Import scikit-learn metrics module for accuracy calculation
from sklearn import metrics

#Load dataset
cancer = datasets.load_breast_cancer()

# print the names of the 13 features
print("Features: ", cancer.feature_names)

# print the label type of cancer('malignant' 'benign')
print("Labels: ", cancer.target_names)

# print data(feature)shape
cancer.data.shape

# print the cancer data features (top 5 records)
print(cancer.data[0:5])

# print the cancer labels (0:malignant, 1:benign)
print(cancer.target)

# Split dataset into training set and test set
X_train, X_test, y_train, y_test = train_test_split(cancer.data, cancer.target,
test_size=0.3,random_state=109) # 70% training and 30% test
```

```
#Create a svm Classifier
```

```
clf = svm.SVC(kernel='linear') # Linear Kernel
```

```
#Train the model using the training sets
```

```
clf.fit(X_train, y_train)
```

```
#Predict the response for test dataset
```

```
y_pred = clf.predict(X_test)
```

```
# Model Accuracy: how often is the classifier correct?
```

```
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

```
# Model Precision: what percentage of positive tuples are labeled as such?
```

```
print("Precision:",metrics.precision_score(y_test, y_pred))
```

```
# Model Recall: what percentage of positive tuples are labelled as such?
```

```
print("Recall:",metrics.recall_score(y_test, y_pred))
```

Output:

```
===== RESTART: E:/MSC/Sem3/AAI Prac/Pranay/P9.py =====
Features: ['mean radius' 'mean texture' 'mean perimeter' 'mean area'
'mean smoothness' 'mean compactness' 'mean concavity'
'mean concave points' 'mean symmetry' 'mean fractal dimension'
'radius error' 'texture error' 'perimeter error' 'area error'
'smoothness error' 'compactness error' 'concavity error'
'concave points error' 'symmetry error' 'fractal dimension error'
'worst radius' 'worst texture' 'worst perimeter' 'worst area'
'worst smoothness' 'worst compactness' 'worst concavity'
'worst concave points' 'worst symmetry' 'worst fractal dimension']
Labels: ['malignant' 'benign']
[[1.799e+01 1.038e+01 1.228e+02 1.001e+03 1.184e-01 2.776e-01 3.001e-01
 1.471e-01 2.419e-01 7.871e-02 1.095e+00 9.053e-01 8.589e+00 1.534e+02
 6.399e-03 4.904e-02 5.373e-02 1.587e-02 3.003e-02 6.193e-03 2.538e+01
 1.733e+01 1.846e+02 2.019e+03 1.622e-01 6.656e-01 7.119e-01 2.654e-01
 4.601e-01 1.189e-01]
[2.057e+01 1.777e+01 1.329e+02 1.326e+03 8.474e-02 7.864e-02 8.690e-02
 7.017e-02 1.812e-01 5.667e-02 5.435e-01 7.339e-01 3.398e+00 7.408e+01
 5.225e-03 1.308e-02 1.860e-02 1.340e-02 1.389e-02 3.532e-03 2.499e+01
 2.341e+01 1.588e+02 1.956e+03 1.238e-01 1.866e-01 2.416e-01 1.860e-01
 2.750e-01 8.902e-02]
[1.969e+01 2.125e+01 1.300e+02 1.203e+03 1.096e-01 1.599e-01 1.974e-01
 1.279e-01 2.069e-01 5.999e-02 7.456e-01 7.869e-01 4.585e+00 9.403e+01
 6.150e-03 4.006e-02 3.832e-02 2.058e-02 2.250e-02 4.571e-03 2.357e+01
 2.553e+01 1.525e+02 1.709e+03 1.444e-01 4.245e-01 4.504e-01 2.430e-01
 3.613e-01 8.758e-02]
[1.142e+01 2.038e+01 7.758e+01 3.861e+02 1.425e-01 2.839e-01 2.414e-01
 1.052e-01 2.597e-01 9.744e-02 4.956e-01 1.156e+00 3.445e+00 2.723e+01
 9.110e-03 7.458e-02 5.661e-02 1.867e-02 5.963e-02 9.208e-03 1.491e+01
 2.650e+01 9.887e+01 5.677e+02 2.098e-01 8.663e-01 6.869e-01 2.575e-01
 6.638e-01 1.730e-01]
[2.029e+01 1.434e+01 1.351e+02 1.297e+03 1.003e-01 1.328e-01 1.980e-01
 1.043e-01 1.809e-01 5.883e-02 7.572e-01 7.813e-01 5.438e+00 9.444e+01
 1.149e-02 2.461e-02 5.688e-02 1.885e-02 1.756e-02 5.115e-03 2.254e+01
 1.667e+01 1.522e+02 1.575e+03 1.374e-01 2.050e-01 4.000e-01 1.625e-01
 2.364e-01 7.678e-02]]

[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0
 1 0 0 0 0 0 0 0 0 0 1 0 1 1 1 1 1 0 0 1 0 0 1 1 1 1 0 1 0 0 1 1 1 1 0 1 0 0
 1 0 1 0 0 1 1 1 0 0 1 0 0 0 1 1 1 0 1 1 0 0 1 1 1 0 0 1 1 1 1 0 1 1 0 1 1
 1 1 1 1 1 1 0 0 0 1 0 0 1 1 1 0 0 1 0 1 0 0 1 0 0 1 1 0 1 1 0 1 1 1 1 0 1
 1 1 1 1 1 1 1 1 0 0 1 1 1 1 0 0 1 0 1 1 0 0 1 1 0 0 1 1 1 1 0 1 1 0 0 0 1 0
 1 0 1 1 1 0 1 1 0 0 1 0 0 0 0 1 0 0 0 1 0 1 0 1 1 0 1 1 0 1 0 0 0 0 1 1 0 0 1 1
 1 0 1 1 1 1 1 0 0 1 1 0 1 1 0 0 1 0 1 1 1 1 1 0 1 1 1 1 1 0 1 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 1 1 1 1 1 1 0 1 0 1 1 0 1 1 0 1 0 0 1 1 1 1 1 1 1 1 1 1 1 1
 1 0 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 0 1 0 1 1 1 1 0 0 0 0 1 1
 1 1 0 1 0 1 0 1 1 1 0 1 1 1 1 1 1 1 1 0 0 0 1 1 1 1 1 1 1 1 1 1 1 0 0 1 0 0
 0 1 0 0 1 1 1 1 1 0 1 1 1 1 1 0 1 1 1 0 1 1 0 0 1 1 1 1 1 1 0 1 1 1 1 1 1 1
 1 0 1 1 1 1 1 0 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 0 1 0 0 1 0 1 1 1 1 1 0 1 1
 0 1 0 1 1 0 1 0 1 1 1 1 1 1 1 1 0 0 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 0 1
 1 1 1 1 1 1 0 1 0 1 1 1 1 1 1 1 1 0 0 1 0 1 0 1 1 1 1 1 0 1 1 0 1 0 1 0 0
 1 1 1 0 1 1 1 1 1 1 1 1 1 1 0 1 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
 1 1 1 1 1 1 1 1 0 0 0 0 0 0 1]
Accuracy: 0.9649122807017544
Precision: 0.9811320754716981
Recall: 0.9629629629629629
>>
```

Practical 10: Simulate artificial neural network model with both feedforward and backpropagation approach.

```
import numpy as np

X = np.array([[2, 9], [1, 5], [3, 6]], dtype=float) # two inputs [sleep,study]
y = np.array([92, 86, 89], dtype=float) # one output [Expected % in Exams]
X = X / np.amax(X, axis=0) # maximum of X array longitudinally
y = y / 100

# Sigmoid Function
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

# Derivative of Sigmoid Function
def derivatives_sigmoid(x):
    return x * (1 - x)

# Variable initialization
epoch = 5000 # Setting training iterations
lr = 0.1 # Setting learning rate
inputlayer_neurons = 2 # number of features in data set
hiddenlayer_neurons = 3 # number of hidden layers neurons
output_neurons = 1 # number of neurons at output layer

# weight and bias initialization
wh = np.random.uniform(size=(inputlayer_neurons, hiddenlayer_neurons)) # weight of the link from
input node to hidden node
bh = np.random.uniform(size=(1, hiddenlayer_neurons)) # bias of the link from input node to
hidden node
wout = np.random.uniform(size=(hiddenlayer_neurons, output_neurons)) # weight of the link from
hidden node to output node
bout = np.random.uniform(size=(1, output_neurons)) # bias of the link from hidden node to output
node
```



```
# draws a random range of numbers uniformly of dim x*y
```

```
for i in range(epoch):  
    # Forward Propagation  
    hinp1 = np.dot(X, wh)  
    hinp = hinp1 + bh  
    hlayer_act = sigmoid(hinp)  
    outinp1 = np.dot(hlayer_act, wout)  
    outinp = outinp1 + bout  
    output = sigmoid(outinp)  
    # Backpropagation  
    EO = y - output  
    outgrad = derivatives_sigmoid(output)  
    d_output = EO * outgrad  
    EH = d_output.dot(wout.T)  
    # how much hidden layer weights contributed to error  
    hiddengrad = derivatives_sigmoid(hlayer_act)  
    d_hiddenlayer = EH * hiddengrad  
    # dotproduct of nextlayererror and currentlayerop  
    wout += hlayer_act.T.dot(d_output) * lr  
    wh += X.T.dot(d_hiddenlayer) * lr  
    print("Input: \n" + str(X))  
    print("Actual Output: \n" + str(y))  
    print("Predicted Output: \n", output)
```

Output:

```
===== RESTART: E:/MSC/Sem3/AAI Prac/Pranay/P10.py =====  
Input:  
[[0.66666667 1.          ]  
 [0.33333333 0.55555556]  
 [1.          0.66666667]]  
Actual Output:  
[[0.92]  
 [0.86]  
 [0.89]]  
Predicted Output:  
[[0.69044378]  
 [0.67208465]  
 [0.68799528]]  
>>|
```


6.1	Implement Bayesian Linear Regression to explore prior and posterior distribution	50
6.2	Implement Gaussian Mixture Models for density estimation and unsupervised clustering	52
Practical 7		54
7.1	Implement cross-validation techniques (k-fold, stratified, etc.) for robust model evaluation	54
7.2	Systematically explore combinations of hyperparameters to optimize model performance.(use grid and randomized search)	55
Practical 8		57
8.1	Implement Bayesian Learning using inferences.....	57
Practical 9		58
9.1	Set up a generator network to produce samples and a discriminator network to distinguish between real and generated data. (Use a simple small dataset).....	58
Practical 10		62
10.1	Develop an API to deploy your model and perform predictions	62

Practical 1

1.1 Load a CSV dataset. Handle missing values, inconsistent formatting, and outliers.

```
import seaborn as sns
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
plt.rcParams["figure.figsize"] = [8,6]
```

```
sns.set(style="darkgrid")
```

```
df = sns.load_dataset('titanic')
```

```
print(df.head())
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class
0	0	3	male	22.0	1	0	7.2500	S	Third
1	1	1	female	38.0	1	0	71.2833	C	First
2	1	3	female	26.0	0	0	7.9250	S	Third
3	1	1	female	35.0	1	0	53.1000	S	First
4	0	3	male	35.0	0	0	8.0500	S	Third

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True

```
print(df.isnull().sum())
```

survived	0
pclass	0
sex	0
age	177
sibsp	0
parch	0
fare	0
embarked	2
class	0
who	0
adult_male	0

```

alone          0
dtype: int64

df = df[["age", "embarked"]]
print(df.head())

   age embarked
0  22.0        S
1  38.0        C
2  26.0        S
3  35.0        S
4  35.0        S

df.loc[:, 'age'] = df.age.fillna(df.age.median())
df = df.dropna(subset=["embarked"])

print(df.head(20))

   age embarked
0  22.0        S
1  38.0        C
2  26.0        S
3  35.0        S
4  35.0        S
5  28.0        Q
6  54.0        S
7   2.0        S
8  27.0        S
9  14.0        C
10  4.0        S
11 58.0        S
12 20.0        S
13 39.0        S
14 14.0        S
15 55.0        S
16  2.0        Q
17 28.0        S
18 31.0        S
19 28.0        C


df.loc[:, 'embarked'] = df.embarked.str.upper()
print(df.embarked.unique())

['S' 'C' 'Q']

sns.boxplot(data=df.age)

<Axes: ylabel='age'>

```



```
Q1 = df.age.quantile(0.25)
Q3 = df.age.quantile(0.75)
IQR = Q3 - Q1
```

```
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
```


```
df = df[(df.age >= lower_bound) & (df.age <= upper_bound)]
```

```
print(df.head())
```

	age	embarked
0	22.0	S
1	38.0	C
2	26.0	S
3	35.0	S
4	35.0	S

```
sns.boxplot(data=df.age)
```

```
<Axes: ylabel='age'>
```



1.2 Load a dataset, calculate descriptive summary statistics, create visualizations using different graphs, and identify potential features and target variables

```
import pandas as pd
import numpy as np
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
plt.rcParams["figure.figsize"] = [8,6]
sns.set(style="darkgrid")
```

```
df = sns.load_dataset('iris')
print(df.head())
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa


```
summary_statistics = df.describe()
print(summary_statistics)
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

#Univariate Visualizations


```
df.hist()
```

```
array([[<Axes: title={'center': 'sepal_length'}>,
        <Axes: title={'center': 'sepal_width'}>],
       [<Axes: title={'center': 'petal_length'}>,
        <Axes: title={'center': 'petal_width'}>]], dtype=object)
```





```
sns.boxplot(data=df)
```

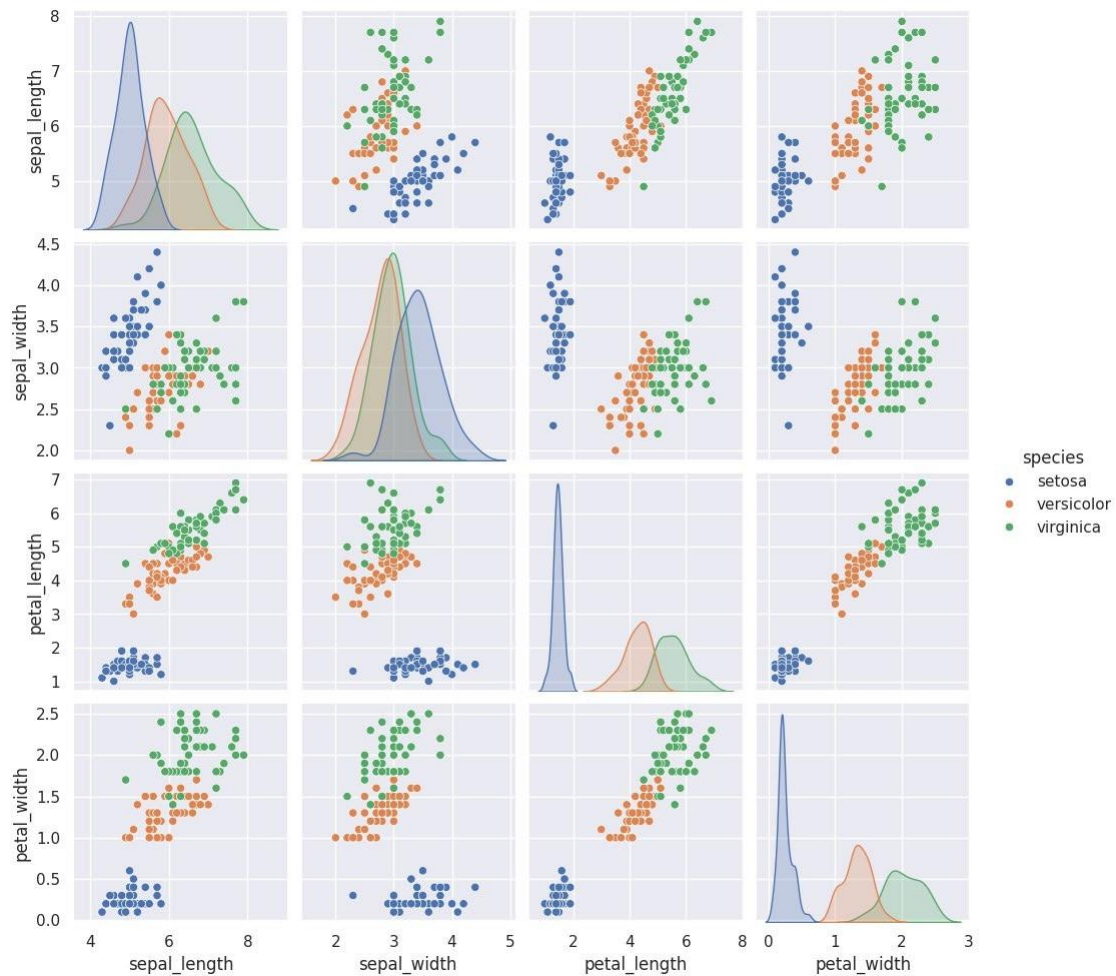
```
<Axes: >
```



```
sns.scatterplot(data=df, x='sepal_length', y='sepal_width', hue='species')  
<Axes: xlabel='sepal_length', ylabel='sepal_width'>
```



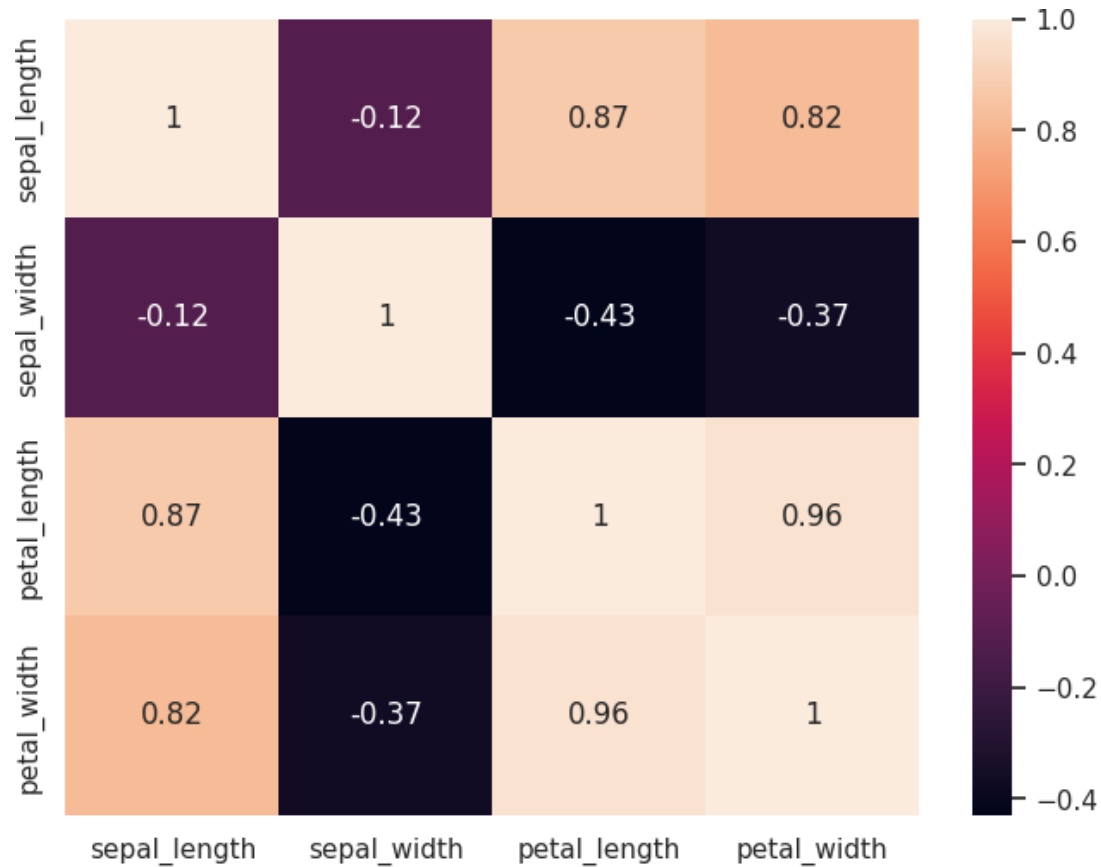
```
sns.pairplot(df, hue="species")  
<seaborn.axisgrid.PairGrid at 0x7fc52c528610>
```



#Correlation

```
numeric_df = df.select_dtypes(include=[np.number])
correlation_matrix = numeric_df.corr()
sns.heatmap(correlation_matrix, annot=True)
```

<Axes: >



```
potential_features = df.select_dtypes(include=[np.number]).columns.tolist()
)
```

```
target_variable = 'species'
```

```
print("Potential Features: ", potential_features)
print("Target Variable: ", target_variable)
```

```
Potential Features: ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']
Target Variable: species
```

1.3 Create or Explore datasets to use all pre-processing routines like label encoding, scaling, and binarization.

```
import seaborn as sns
```

```
import pandas as pd
```

```
from sklearn.preprocessing import LabelEncoder from  
sklearn.preprocessing import StandardScaler from  
sklearn.preprocessing import Binarizer
```

```
df = sns.load_dataset('titanic')  
print(df.head())
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class
0	0	3	male	22.0	1	0	7.2500	S	Third
1	1	1	female	38.0	1	0	71.2833	C	First
2	1	3	female	26.0	0	0	7.9250	S	Third
3	1	1	female	35.0	1	0	53.1000	S	First
4	0	3	male	35.0	0	0	8.0500	S	Third

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True

```
label_encoder = LabelEncoder()  
df['sex'] = label_encoder.fit_transform(df['sex'])  
print(df.sex.head())
```

```
0    1  
1    0  
2    0  
3    0  
4    1  
Name: sex, dtype: int64
```

```
scaler = StandardScaler()  
df[['age', 'fare']] = scaler.fit_transform(df[['age', 'fare']])  
print(df[['age', 'fare']].head())
```

	age	fare
0	-0.530377	-0.502445
1	0.571831	0.786845
2	-0.254825	-0.488854

```
binarizer = Binarizer(threshold=0)
df[['fare']] = binarizer.fit_transform(df[['fare']])
print(df.fare.head())
```

```
0    0.0
1    1.0
2    0.0
3    1.0
4    0.0
```

```
Name: fare, dtype: float64
```

Practical 2

2.1 Implement and demonstrate the FIND-S algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from a CSV file and generate the final specific hypothesis. (Create your dataset)

```
import pandas as pd

data = {
    'Material': ['Plastic', 'Metal', 'Glass', 'Metal'],
    'Color': ['White', 'Silver', 'Green', 'Grey'],
    'Size': ['Small', 'Large', 'Small', 'Large'],
    'Recyclable': ['Yes', 'Yes', 'Yes', 'No'],
    'E-Waste': ['No', 'Yes', 'No', 'Yes']
}

df = pd.DataFrame(data)
df.to_csv('training_data.csv', index=False)

data = pd.read_csv('training_data.csv')
X = data.iloc[:, :-1]
y = data.iloc[:, -1]

hypothesis = ['0'] * X.shape[1]

for i in range(len(X)):
    if y[i] == 'Yes':
        for j in range(X.shape[1]):
            if hypothesis[j] == '0':
                hypothesis[j] = X.iloc[i, j]
            elif hypothesis[j] != X.iloc[i, j]:
                hypothesis[j] = '?'

print(hypothesis)

['Metal', '?', 'Large', '?']
```


Practical 3

3.1 Simple Linear Regression Fit a linear regression model on a dataset. Interpret coefficients, make predictions, and evaluate performance using metrics like R-squared and MSE

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

np.random.seed(42)
X = 2.5 * np.random.rand(100, 1)
y = 5 + 2 * X + np.random.randn(100, 1)

data = pd.DataFrame({'Feature': X.flatten(), 'Target': y.flatten()})
print(data.head())
```

	Feature	Target
0	0.936350	6.959748
1	2.376786	9.454564
2	1.829985	8.751730
3	1.496646	6.005724
4	0.390047	5.560421

```
plt.figure(figsize=(8, 6))
sns.scatterplot(x='Feature', y='Target', data=data)
plt.title('Feature vs Target')
plt.show()
```



```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
random_state=42)
```

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

```
print(f"Intercept: {model.intercept_[0]:.2f}")
```

```
print(f"Coefficient: {model.coef_[0][0]:.2f}")
```

```
Intercept: 5.14
```

```
Coefficient: 1.84
```

```
y_pred = model.predict(X_test)
```

```
pred_df = pd.DataFrame({'Actual': y_test.flatten(), 'Predicted': y_pred.fl
atten()})
```

```
print(pred_df.head())
```

	Actual	Predicted
0	5.974345	5.435196
1	8.970661	9.257909
2	7.624273	8.694195


```
3  7.403224  8.189620
4  7.084932  6.332951
```

```
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error (MSE): {mse:.2f}")
```

```
r2 = r2_score(y_test, y_pred)
print(f"R-squared: {r2:.2f}")
```

Mean Squared Error (MSE): 0.65
R-squared: 0.73

```
plt.figure(figsize=(8, 6))
sns.scatterplot(x='Feature', y='Target', data=data, label='Actual Data')
plt.plot(X_test, y_pred, color='red', label='Regression Line')
plt.title('Linear Regression Fit')
plt.legend()
plt.show()
```



3.2 Multiple Linear Regression Extend linear regression to multiple features. Handle feature selection and potential multicollinearity.

```
import seaborn as sns
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt

import statsmodels.api as sm

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from statsmodels.stats.outliers_influence import variance_inflation_factor


df = sns.load_dataset('diamonds')

print("Missing values in the dataset:")
print(df.isnull().sum())

df = pd.get_dummies(df, drop_first=True)

Missing values in the dataset:
carat      0
cut        0
color      0
clarity     0
depth      0
table      0
price      0
x          0
y          0
z          0
dtype: int64

plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title("Correlation Heatmap")
plt.show()
```



```
X = df[['carat', 'depth', 'table', 'x', 'y', 'z',
        'cut_Premium', 'cut_Good', 'cut_Very Good',
        'color_E', 'color_F', 'clarity_VS2', 'clarity_VS1']]
y = df['price']

y = y.astype(float)

X_with_constant = sm.add_constant(X)

X_with_constant = X_with_constant.astype(int)

vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X_with_constant.values, i+1) for i
              in range(len(X.columns))]
print(vif)
```

	Features	VIF
0	carat	3.613538
1	depth	1.211270

```

2         table    1.530672
3         x      19.224267
4         y      15.677513
5         z       5.789510
6     cut_Premium  1.548643
7         cut_Good  1.295429
8     cut_Very Good  1.346362
9         color_E   1.079848
10        color_F   1.060367
11    clarity_VS2   1.051655
12    clarity_VS1   1.031049

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, r
andom_state=42)

```

```

print("Data types of X_train:")
print(X_train.dtypes)
print("Data type of y_train:", y_train.dtype)

```

Data types of X_train:

```

carat          float64
depth          float64
table          float64
x              float64
y              float64
z              float64
cut_Premium    bool
cut_Good       bool
cut_Very Good  bool
color_E        bool
color_F        bool
clarity_VS2    bool
clarity_VS1    bool
dtype: object

```

Data type of y_train: float64

```

X_train = X_train.astype(float)
y_train = y_train.astype(float)

```

```

model = LinearRegression()
model.fit(X_train, y_train)

```

```

print("Intercept:", model.intercept_)
print("Coefficients:", model.coef_)

```

```

Intercept: 17520.480548853404
Coefficients: [ 1.06799558e+04 -1.74848962e+02 -8.87411825e+01 -1.17618393
e+03
 3.03543071e+01  8.16330490e+00 -3.94552416e+01 -1.98436785e+02

```

```
-1.87877044e+01 4.36586414e+02 4.74099129e+02 1.02842658e+03  
6.62232418e+02]
```


```
y_pred = model.predict(X_test)  
mse = mean_squared_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)
```

```
print("Mean Squared Error:", mse)  
print("R-squared:", r2)
```


```
residuals = y_test - y_pred
```

```
Mean Squared Error: 2018911.748442661  
R-squared: 0.8705490836162249
```

```
plt.figure(figsize=(10, 6))  
sns.histplot(residuals, kde=True)  
plt.title("Residuals Distribution")  
plt.show()
```



```
plt.scatter(y_pred, residuals)  
plt.axhline(y=0, color='r', linestyle='--')  
plt.xlabel("Predicted Values")  
plt.ylabel("Residuals")  
plt.title("Residuals vs Predicted Values")  
plt.show()
```



```
X_train_sm = sm.add_constant(X_train)
ols_model = sm.OLS(y_train, X_train_sm).fit()
print(ols_model.summary())
```

OLS Regression Results

```
=====
=====
====
Dep. Variable:          price    R-squared:                0
.870
Model:                  OLS      Adj. R-squared:            0
.869
Method:                 Least Squares    F-statistic:            1.935
e+04
Date:                   Thu, 24 Oct 2024    Prob (F-statistic):
0.00
Time:                   07:44:35    Log-Likelihood:         -3.2835
e+05
No. Observations:       37758    AIC:                    6.567
e+05
Df Residuals:           37744    BIC:                    6.569
e+05
Df Model:                13
Covariance Type:        nonrobust
=====
=====
```


	coef	std err	t	P> t	[0.025	0.975]

const	1.752e+04	537.976	32.567	0.000	1.65e+04	1
.86e+04						
carat	1.068e+04	72.935	146.431	0.000	1.05e+04	1
.08e+04						
depth	-174.8490	6.336	-27.595	0.000	-187.268	-
162.430						
table	-88.7412	4.144	-21.414	0.000	-96.864	
-80.619						
x	-1176.1839	46.936	-25.059	0.000	-1268.179	-1
084.188						
y	30.3543	27.897	1.088	0.277	-24.325	
85.034						
z	8.1633	43.889	0.186	0.852	-77.861	
94.188						
cut_Premium	-39.4552	21.233	-1.858	0.063	-81.073	
2.163						
cut_Good	-198.4368	29.625	-6.698	0.000	-256.502	-
140.372						
cut_Very Good	-18.7877	20.712	-0.907	0.364	-59.383	
21.808						
color_E	436.5864	20.108	21.712	0.000	397.174	
475.999						
color_F	474.0991	20.084	23.606	0.000	434.735	
513.463						
clarity_VVS2	1028.4266	26.287	39.123	0.000	976.903	1
079.950						
clarity_VS1	662.2324	21.075	31.423	0.000	620.925	
703.540						

```

=====
====
Omnibus:                9105.593    Durbin-Watson:                1
.992
Prob(Omnibus):          0.000    Jarque-Bera (JB):            327369
.326
Skew:                   0.453    Prob(JB):
0.00
Kurtosis:               17.397    Cond. No.                    6.14
e+03
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

3.3 Regularized Linear Models (Ridge, Lasso, ElasticNet) Implement regression variants like LASSO and Ridge on any generated dataset.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.datasets import make_regression
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge, Lasso, ElasticNet
from sklearn.metrics import mean_squared_error

X, y, coef = make_regression(n_samples=100, n_features=10, noise=0.1, coef=True, random_state=42)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)



ridge_model = Ridge(alpha=1.0)
ridge_model.fit(X_train, y_train)
ridge_pred = ridge_model.predict(X_test)

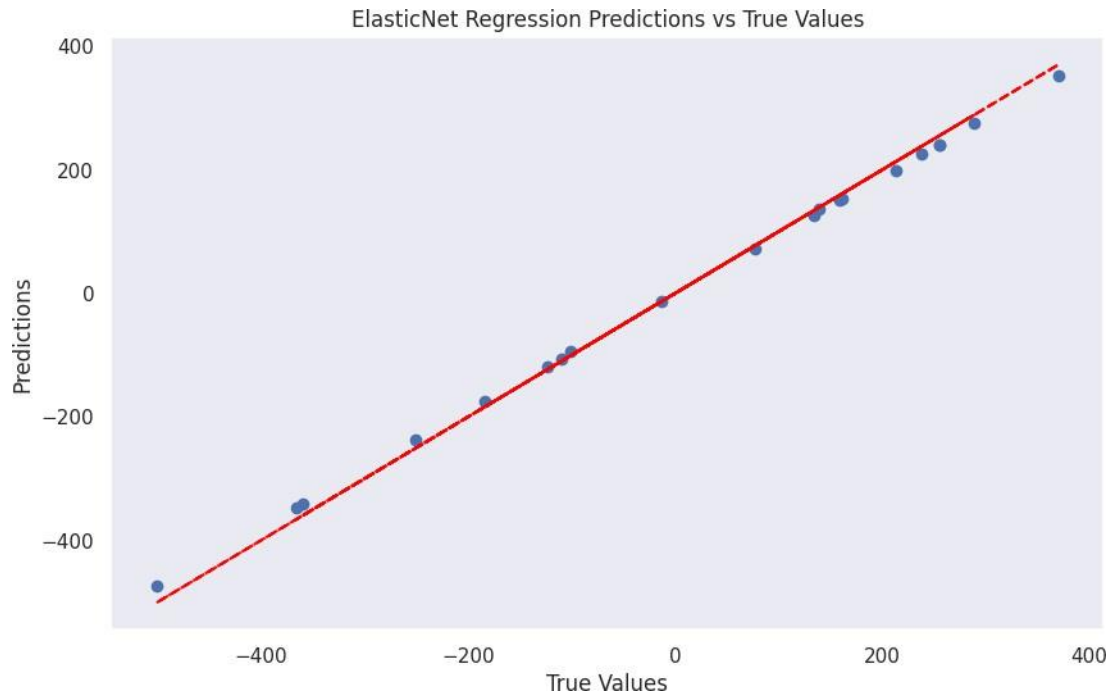
lasso_model = Lasso(alpha=0.1)
lasso_model.fit(X_train, y_train)
lasso_pred = lasso_model.predict(X_test)

elastic_model = ElasticNet(alpha=0.1, l1_ratio=0.5)
elastic_model.fit(X_train, y_train)
elastic_pred = elastic_model.predict(X_test)

def plot_results(y_test, predictions, model_name):
    plt.figure(figsize=(10, 6))
    plt.scatter(y_test, predictions)
    plt.plot(y_test, y_test, color='red', linestyle='--') # y=x line
    plt.title(f'{model_name} Predictions vs True Values')
    plt.xlabel('True Values')
    plt.ylabel('Predictions')
    plt.grid()
    plt.show()

plot_results(y_test, ridge_pred, 'Ridge Regression')
plot_results(y_test, lasso_pred, 'Lasso Regression')
plot_results(y_test, elastic_pred, 'ElasticNet Regression')
```





```
print("Mean Squared Error (MSE):")
print(f"Ridge: {mean_squared_error(y_test, ridge_pred):.2f}")
print(f"Lasso: {mean_squared_error(y_test, lasso_pred):.2f}")
print(f"ElasticNet: {mean_squared_error(y_test, elastic_pred):.2f}")
```

Mean Squared Error (MSE):
Ridge: 11.84
Lasso: 0.18
ElasticNet: 176.03

Practical 4

4.1 Logistic Regression Perform binary classification using logistic regression. Calculate accuracy, precision, recall, and understand the ROC curve.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score,
roc_curve, auc
from sklearn.datasets import make_classification

X, y = make_classification(n_samples=1000, n_features=10, n_classes=2, random_state=42)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
logistic_reg_model = LogisticRegression()
logistic_reg_model.fit(X_train, y_train)
y_pred = logistic_reg_model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)


print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")

Accuracy: 0.85
Precision: 0.89
Recall: 0.82

y_prob = logistic_reg_model.predict_proba(X_test)[:, 1]
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)
plt.figure()
plt.plot(fpr, tpr, color='blue', lw=2, label=f"ROC Curve (AUC = {roc_auc:.2f})")
plt.plot([0, 1], [0, 1], color='gray', linestyle='--') # Diagonal line for random classifier

plt.xlim([0.0, 1.0])
```

```
plt.legend(loc='lower right')  
plt.show()
```



4.2 Implement and demonstrate k-nearest Neighbor algorithm. Read the training data from a .CSV file and build the model to classify a test sample. Print both correct and wrong predictions.

```
from sklearn import datasets
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report

iris = datasets.load_iris()
X = iris.data
y = iris.target

df = pd.DataFrame(data=X, columns=iris.feature_names)
df['target'] = y

print(df.head())
```

```
sepal.length(cm) sepal.width(cm) petal.length(cm) petal.width(cm) target
0  5.1  3.5  1.4  0.
2  4.9  3.0  1.4  0.
2  4.7  3.2  1.3  0.
3  4.6  3.1  1.5  0.
4  5.0  3.6  1.4  0.
2  4.4  2.9  1.4  0.
```


```

X = df.iloc[:, :-1]
y = df.iloc[:, -1]
y = y.astype('category').cat.codes

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
andom_state=42)
k = 3
knn = KNeighborsClassifier(n_neighbors=k)
knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

for i in range(len(y_test)):
    print(f'Predicted: {iris.target_names[y_pred[i]]}, Actual: {iris.targe
t_names[y_test.iloc[i]]}')

Predicted: versicolor, Actual: versicolor
Predicted: setosa, Actual: setosa
Predicted: virginica, Actual: virginica
Predicted: versicolor, Actual: versicolor
Predicted: versicolor, Actual: versicolor
Predicted: setosa, Actual: setosa
Predicted: versicolor, Actual: versicolor
Predicted: virginica, Actual: virginica
Predicted: versicolor, Actual: versicolor
Predicted: versicolor, Actual: versicolor
Predicted: virginica, Actual: virginica
Predicted: setosa, Actual: setosa
Predicted: setosa, Actual: setosa
Predicted: setosa, Actual: setosa
Predicted: setosa, Actual: setosa
Predicted: versicolor, Actual: versicolor
Predicted: virginica, Actual: virginica
Predicted: versicolor, Actual: versicolor
Predicted: versicolor, Actual: versicolor
Predicted: virginica, Actual: virginica
Predicted: setosa, Actual: setosa
Predicted: virginica, Actual: virginica
Predicted: setosa, Actual: setosa
Predicted: virginica, Actual: virginica
Predicted: virginica, Actual: virginica
Predicted: virginica, Actual: virginica
Predicted: virginica, Actual: virginica
Predicted: virginica, Actual: virginica
Predicted: setosa, Actual: setosa
Predicted: setosa, Actual: setosa

accuracy = accuracy_score(y_test, y_pred)

```



```
print(classification_report(y_test, y_pred, target_names=iris.target_names
))
```

Accuracy: 100.00%

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	1.00	1.00	1.00	9
virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

4.3 Build a decision tree classifier or regressor. Control hyperparameters like tree depth to avoid overfitting. Visualize the tree.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.datasets import load_iris, fetch_california_housing
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, plot_tree
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, mean_squared_error

iris = load_iris()
X = iris.data
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

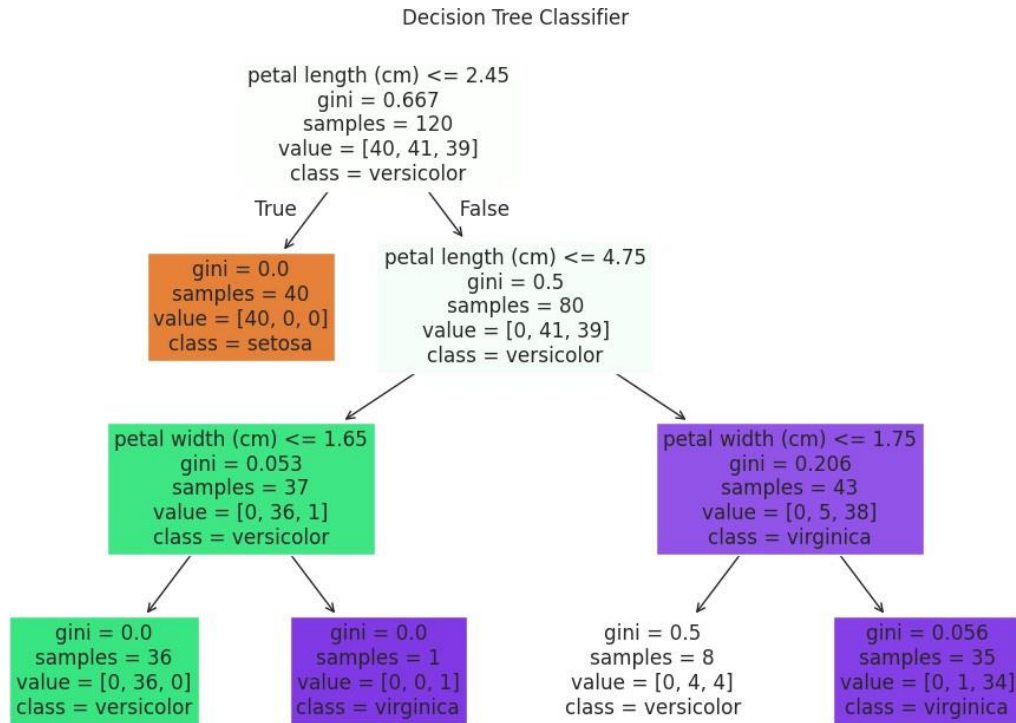
clf = DecisionTreeClassifier(max_depth=3, random_state=42)
clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')

Accuracy: 1.00

plt.figure(figsize=(12,8))
plot_tree(clf, filled=True, feature_names=iris.feature_names, class_names=iris.target_names)
plt.title("Decision Tree Classifier")
plt.show()
```



```

housing = fetch_california_housing()
X_housing = housing.data
y_housing = housing.target

X_train_housing, X_test_housing, y_train_housing, y_test_housing = train_t
est_split(X_housing, y_housing, test_size=0.2, random_state=42)

reg = DecisionTreeRegressor(max_depth=3, random_state=42)
reg.fit(X_train_housing, y_train_housing)

y_pred_housing = reg.predict(X_test_housing)

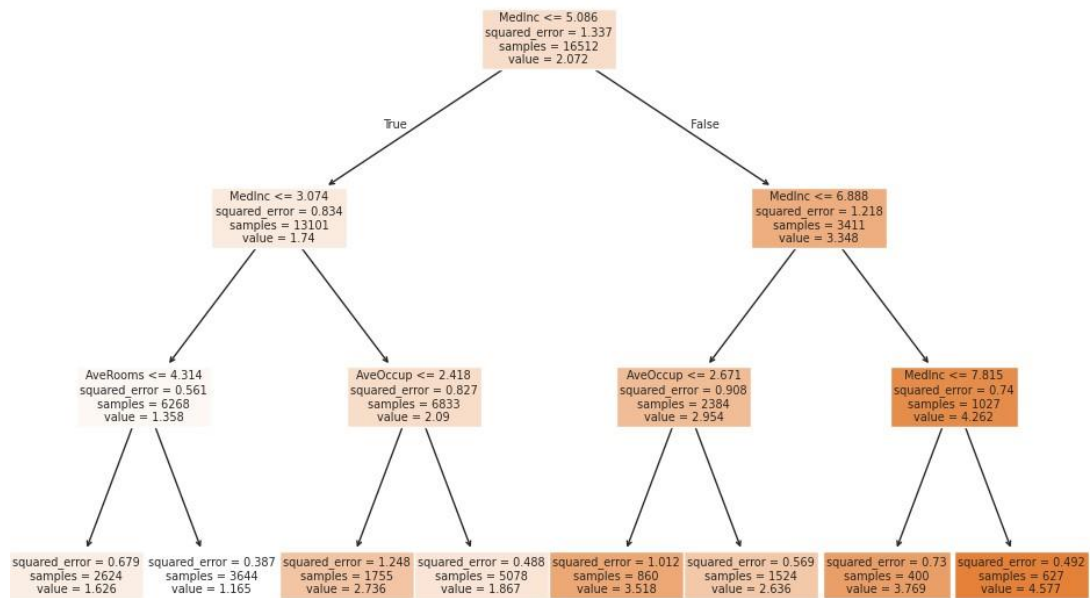
mse = mean_squared_error(y_test_housing, y_pred_housing)
print(f'Mean Squared Error: {mse:.2f}')

Mean Squared Error: 0.64

plt.figure(figsize=(12,8))
plot_tree(reg, filled=True, feature_names=housing.feature_names)
plt.title("Decision Tree Regressor")
plt.show()

```

Decision Tree Regressor



4.4 Implement a Support Vector Machine for any relevant dataset.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns

iris = datasets.load_iris()
X = iris.data
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
svm_classifier = SVC(kernel='linear', random_state=42)
svm_classifier.fit(X_train, y_train)

y_pred = svm_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

conf_matrix = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:")
print(conf_matrix)

class_report = classification_report(y_test, y_pred, target_names=iris.target_names)
print("\nClassification Report:")
print(class_report)
```

Accuracy: 1.00

Confusion Matrix:

```
[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]
```


Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

setosa	1.00	1.00	1.00	19
versicolor	1.00	1.00	1.00	13
virginica	1.00	1.00	1.00	13

accuracy				1.00	45
macro avg	1.00	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	1.00	45

```
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=iris.target_names, yticklabels=iris.target_names)
plt.title('Confusion Matrix')
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.show()
```



4.5 Train a random forest ensemble. Experiment with the number of trees and feature sampling. Compare performance to a single decision tree.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import datasets

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns

iris = datasets.load_iris()
X = iris.data
y = iris.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
tree_classifier = DecisionTreeClassifier(random_state=42)
tree_classifier.fit(X_train, y_train)
y_pred_tree = tree_classifier.predict(X_test)
accuracy_tree = accuracy_score(y_test, y_pred_tree)
print(f"Decision Tree Accuracy: {accuracy_tree:.2f}\n")

Decision Tree Accuracy: 1.00

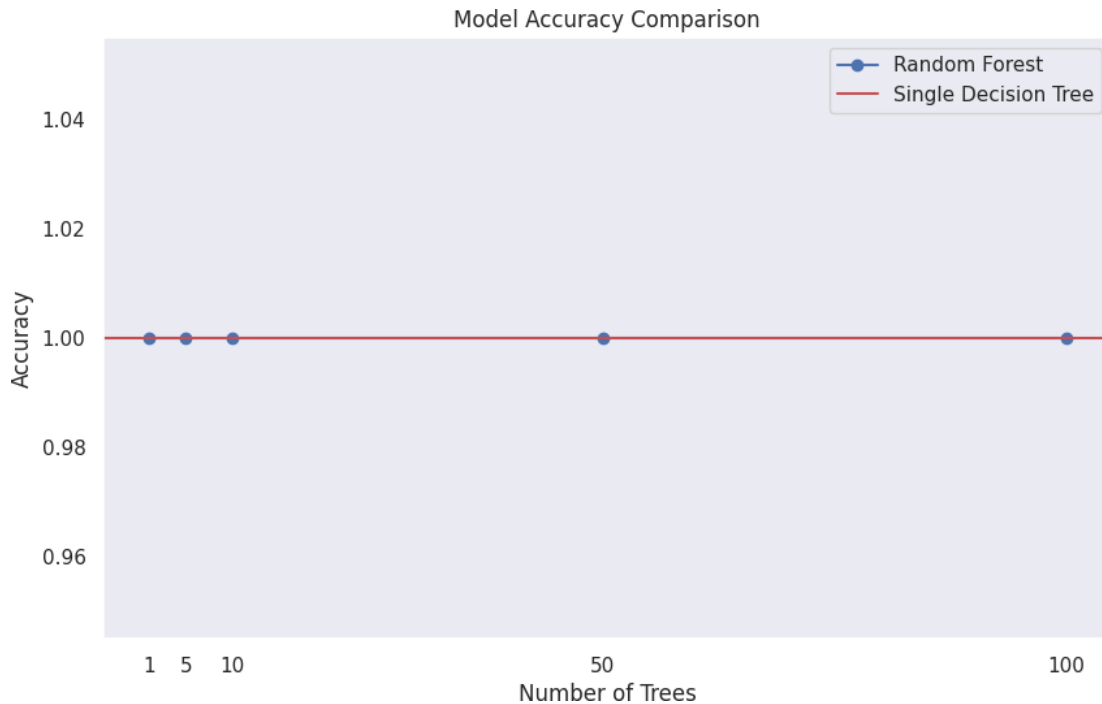
n_trees = [1, 5, 10, 50, 100]
accuracy_forest = []
for n in n_trees:
    forest_classifier = RandomForestClassifier(n_estimators=n, random_state=42)
    forest_classifier.fit(X_train, y_train)
    y_pred_forest = forest_classifier.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred_forest)
    accuracy_forest.append(accuracy)
    print(f"Random Forest with {n} trees Accuracy: {accuracy:.2f}")

Random Forest with 1 trees Accuracy: 1.00
Random Forest with 5 trees Accuracy: 1.00
Random Forest with 10 trees Accuracy: 1.00
Random Forest with 50 trees Accuracy: 1.00
```


Random Forest with 100 trees Accuracy: 1.00

```
plt.figure(figsize=(10, 6))
plt.plot(n_trees, accuracy_forest, marker='o', label='Random Forest')
plt.axhline(y=accuracy_tree, color='r', linestyle='--', label='Single Decision Tree')
plt.title('Model Accuracy Comparison')
```

```
plt.xlabel('Number of Trees')
plt.ylabel('Accuracy')
plt.xticks(n_trees)
plt.legend()
plt.grid()
plt.show()
```



```
best_n = n_trees[np.argmax(accuracy_forest)] # Get the best performing number of trees

best_forest_classifier = RandomForestClassifier(n_estimators=best_n, random_state=42)
best_forest_classifier.fit(X_train, y_train)
y_pred_best_forest = best_forest_classifier.predict(X_test)

conf_matrix = confusion_matrix(y_test, y_pred_best_forest)
print("\nConfusion Matrix for Random Forest (best model):")
print(conf_matrix)


Class_report = classification_report(y_test, y_pred_best_forest, target_names=iris.target_names)
print("\nClassification Report for Random Forest (best model):")
print(class_report)

Confusion Matrix for Random Forest (best model):
[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]
```

Classification Report for Random Forest (best model):

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	19
versicolor	1.00	1.00	1.00	13
virginica	1.00	1.00	1.00	13
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

```
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=iris.target_names, yticklabels=iris.target_names)
plt.title('Confusion Matrix for Random Forest (best model)')
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.show()
```



4.6 Implement a gradient boosting machine (e.g., XGBoost). Tune hyperparameters and explore feature importance.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn import datasets

from sklearn.model_selection import train_test_split, GridSearchCV
from xgboost import XGBClassifier

from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

iris = datasets.load_iris()
X = iris.data
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
xgb_model = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss')
xgb_model.fit(X_train, y_train)

param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [3, 4, 5],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': [0.8, 1.0]
}

grid_search = GridSearchCV(estimator=xgb_model, param_grid=param_grid, scoring='accuracy', cv=3, verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)

best_params = grid_search.best_params_
print("Best parameters from GridSearch:", best_params)

Fitting 3 folds for each of 54 candidates, totalling 162 fits
Best parameters from GridSearch: {'learning_rate': 0.01, 'max_depth': 4, 'n_estimators': 100, 'subsample': 0.8}

best_xgb_model = grid_search.best_estimator_
y_pred = best_xgb_model.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
print(f"\nXGBoost Accuracy: {accuracy:.2f}\n")

conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)
```

```

class_report = classification_report(y_test, y_pred, target_names=iris.target_names)
print("\nClassification Report:")
print(class_report)

```

XGBoost Accuracy: 1.00

Confusion Matrix:

```

[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]

```


Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	19
versicolor	1.00	1.00	1.00	13
virginica	1.00	1.00	1.00	13
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

```

plt.figure(figsize=(10, 6))
sns.barplot(x=best_xgb_model.feature_importances_, y=iris.feature_names)
plt.title('Feature Importance from XGBoost')
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.show()

```



Practical 5

5.1 Implement and demonstrate the working of a Naive Bayesian classifier using a sample data set. Build the model to classify a test sample.

```
import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

iris = load_iris()
X = iris.data
y = iris.target

df = pd.DataFrame(data=X, columns=iris.feature_names)
df['target'] = y

print(df.head())
```

```
sepal length (cm) \ ) \ sepal width (cm) \ ) \ petal length (cm) \ ) \ petal width (cm) \ ) \
0 5.1 5.1 3.5 3.5 1.4 1.4 0. 0.
2 4.9 4.9 3.0 3.0 1.4 1.4 0. 0.
2 4.7 4.7 3.2 3.2 1.3 1.3 0. 0.
2 4.6 4.6 3.1 3.1 1.5 1.5 0. 0.
4 5.0 5.0 3.6 3.6 1.4 1.4 0. 0.
2
```



```

accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

```

```

print(f"Accuracy: {accuracy:.2f}")
print("Confusion Matrix:")
print(conf_matrix)
print("Classification Report:")
print(class_report)

```

Accuracy: 0.98

Confusion Matrix:

```

[[19  0  0]
 [ 0 12  1]
 [ 0  0 13]]

```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	0.92	0.96	13
2	0.93	1.00	0.96	13
accuracy			0.98	45
macro avg	0.98	0.97	0.97	45
weighted avg	0.98	0.98	0.98	45

5.2 Implement Hidden Markov Models using hmmlearn

```
!pip install hmmlearn
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from hmmlearn import hmm
```

```
Requirement already satisfied: hmmlearn in /usr/local/lib/python3.10/dist-packages (0.3.2)
```

```
Requirement already satisfied: numpy>=1.10 in /usr/local/lib/python3.10/dist-packages (from hmmlearn) (1.26.4)
```

```
Requirement already satisfied: scikit-learn!=0.22.0,>=0.16 in /usr/local/lib/python3.10/dist-packages (from hmmlearn) (1.5.2)
```

```
Requirement already satisfied: scipy>=0.19 in /usr/local/lib/python3.10/dist-packages (from hmmlearn) (1.13.1)
```

```
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn!=0.22.0,>=0.16->hmmlearn) (1.4.2)
```

```
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn!=0.22.0,>=0.16->hmmlearn) (3.5.0)
```

```
np.random.seed(42)
```

```
n_samples = 1000
```

```
n_states = 2
```

```
trans_probs = np.array([[0.7, 0.3],  
                        [0.4, 0.6]])
```

```
means = np.array([[1.0], [0.5]])
```

```
covars = np.array([[0.1], [0.2]])
```

```
model = hmm.GaussianHMM(n_components=n_states, covariance_type="diag", n_iter=100)
```

```
model.startprob_ = np.array([0.6, 0.4])
```

```
model.transmat_ = trans_probs
```

```
model.means_ = means
```

```
model.covars_ = covars
```

```
X, Z = model.sample(n_samples)
```

```
plt.figure(figsize=(10, 6))
```

```
plt.plot(X, label='Observed Data', color='grey', alpha=0.5)
```

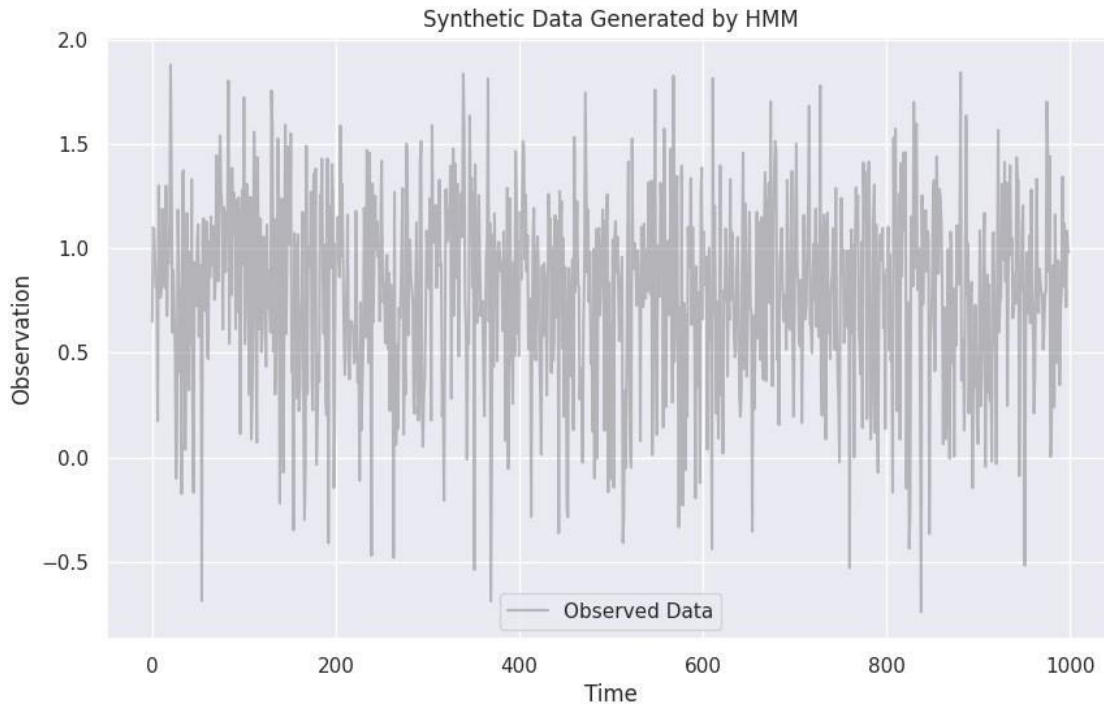
```
plt.title('Synthetic Data Generated by HMM')
```

```
plt.xlabel('Time')
```

```
plt.ylabel('Observation')
```

```
plt.legend()
```

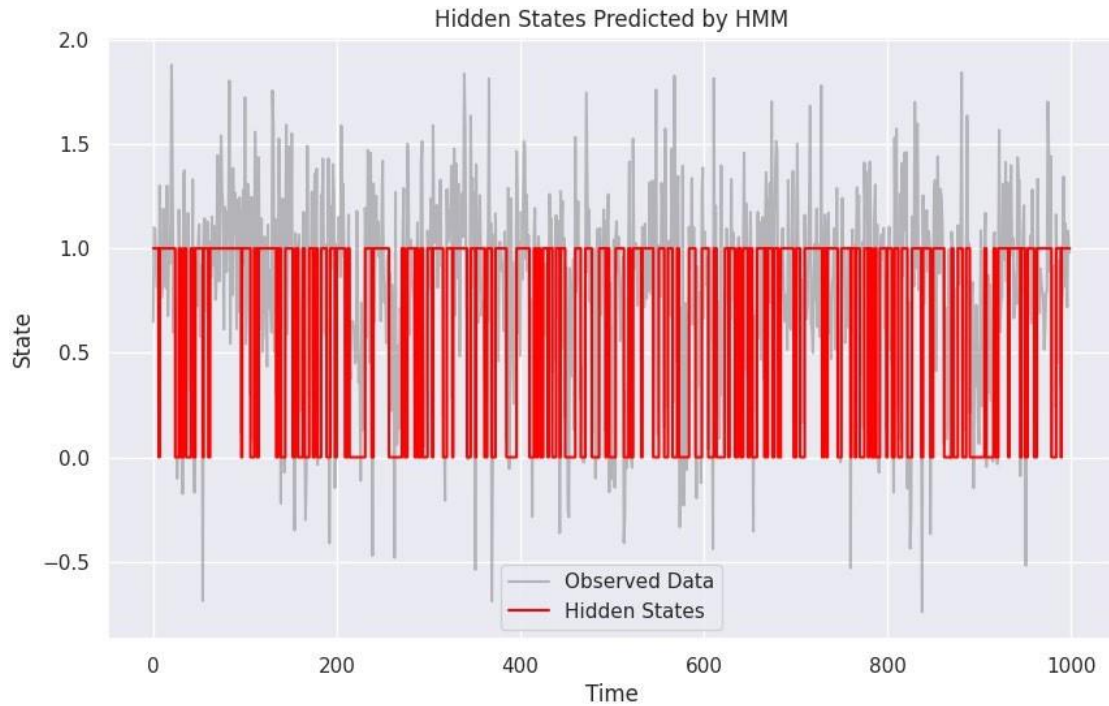
```
plt.show()
```



```
model.fit(X)
hidden_states = model.predict(X)
```

```
WARNING:hmmlearn.base:Even though the 'startprob_' attribute is set, it will
be overwritten during initialization because 'init_params' contains 's'
WARNING:hmmlearn.base:Even though the 'transmat_' attribute is set, it will
be overwritten during initialization because 'init_params' contains 't'
WARNING:hmmlearn.base:Even though the 'means_' attribute is set, it will be
overwritten during initialization because 'init_params' contains 'm'
WARNING:hmmlearn.base:Even though the 'covars_' attribute is set, it will
be overwritten during initialization because 'init_params' contains 'c'
```

```
plt.figure(figsize=(10, 6))
plt.plot(X, label='Observed Data', color='grey', alpha=0.5)
plt.step(range(n_samples), hidden_states, where="post", label='Hidden States',
color='red')
plt.title('Hidden States Predicted by HMM')
plt.xlabel('Time')
plt.ylabel('State')
plt.legend()
plt.show()
```



```
print("Transition matrix:\n", model.transmat_)
print("Means:\n", model.means_)
print("Covariances:\n", model.covars_)
```

Transition matrix:

```
[[0.65865532 0.34134468]
 [0.3121865  0.6878135  ]]
```

Means:

```
[[0.54954006]
 [1.00338912]]
```

Covariances:

```
[[[0.22176075]]

 [[0.09283459]]]
```

Practical 6

6.1 Implement Bayesian Linear Regression to explore prior and posterior distribution.

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```


```
from scipy.stats import multivariate_normal
```

```
np.random.seed(42)
X = np.random.rand(100, 1) * 10
true_beta = np.array([2.0])
y = 2.0 * X.flatten() + np.random.normal(0, 2, size=X.shape[0])
X_b = np.c_[np.ones((X.shape[0], 1)), X]
sigma_0 = 10
sigma_n = 4
sigma_0_inv = 1 / sigma_0
sigma_n_inv = 1 / sigma_n
N = X_b.shape[0]
beta_prior_mean = np.zeros(X_b.shape[1])
beta_prior_cov = sigma_0 * np.eye(X_b.shape[1])
posterior_cov = np.linalg.inv(sigma_n_inv * (X_b.T @ X_b) + sigma_0_inv *
np.eye(X_b.shape[1]))
posterior_mean = posterior_cov @ (sigma_n_inv * (X_b.T @ y))
beta_samples = np.random.multivariate_normal(posterior_mean, posterior_cov
, size=1000)
plt.figure(figsize=(10, 6))
beta_prior_samples = np.random.multivariate_normal(beta_prior_mean, beta_p
rior_cov, size=1000)
```

<Figure size 1000x600 with 0 Axes>

```
plt.subplot(1, 2, 1)
plt.title("Prior Distribution")
plt.scatter(beta_prior_samples[:, 0], beta_prior_samples[:, 1], alpha=0.2)
plt.xlabel("$\\beta_0$")
plt.ylabel("$\\beta_1$")
plt.subplot(1, 2, 2)
plt.title("Posterior Distribution")
plt.scatter(beta_samples[:, 0], beta_samples[:, 1], alpha=0.2, color='orange')
plt.xlabel("$\\beta_0$")
plt.ylabel("$\\beta_1$")
plt.tight_layout()
```

```
plt.show()
```



```
print("Posterior Mean:", posterior_mean)
print("Posterior Covariance:\n", posterior_cov)
```

Posterior Mean: [0.42825291 1.90809351]

Posterior Covariance:

```
[[ 0.1389247 -0.0211579 ]
 [-0.0211579  0.00451795]]
```


6.2 Implement Gaussian Mixture Models for density estimation and unsupervised clustering

```
import numpy as np
import matplotlib.pyplot as plt


from sklearn.mixture import GaussianMixture

np.random.seed(42)
means = [[2, 2], [8, 8], [5, 1]]
covariances = [[[1, 0], [0, 1]], [[1, 0], [0, 1]], [[1, 0], [0, 1]]]
n_samples = 500
data = np.vstack([
    np.random.multivariate_normal(mean, cov, n_samples // len(means))
    for mean, cov in zip(means, covariances)
])

n_components = len(means) # Number of clusters

gmm = GaussianMixture(n_components=n_components, covariance_type='full')
gmm.fit(data)
labels = gmm.predict(data)

plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.scatter(data[:, 0], data[:, 1], c=labels, s=30, cmap='viridis', alpha=0.5)
plt.title('GMM Clustering')
plt.xlabel('X1')
plt.ylabel('X2')
x = np.linspace(-1, 10, 100)
y = np.linspace(-1, 10, 100)
X, Y = np.meshgrid(x, y)
XX = np.column_stack([X.ravel(), Y.ravel()])
logprob = gmm.score_samples(XX)
pdf = np.exp(logprob).reshape(X.shape)
plt.subplot(1, 2, 2)
plt.contourf(X, Y, pdf, levels=20, cmap='Blues', alpha=0.7)
plt.scatter(data[:, 0], data[:, 1], c='red', s=10, alpha=0.5)
plt.title('Density Estimation using GMM')
plt.xlabel('X1')
plt.ylabel('X2')
plt.colorbar(label='Density')
plt.tight_layout()
plt.show()
```



Practical 7

7.1 Implement cross-validation techniques (k-fold, stratified, etc.) for robust model evaluation

```
import numpy as np

from sklearn.model_selection import KFold, StratifiedKFold

from sklearn.metrics import accuracy_score

from sklearn.datasets import load_iris

from sklearn.ensemble import RandomForestClassifier

data = load_iris()
X, y = data.data, data.target
model = RandomForestClassifier()
kf = KFold(n_splits=5, shuffle=True, random_state=42)
kf_scores = []
for train_index, test_index in kf.split(X):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)
    score = accuracy_score(y_test, predictions)
    kf_scores.append(score)
print(f'K-Fold Accuracy: {np.mean(kf_scores):.2f} ± {np.std(kf_scores):.2f}')
```

K-Fold Accuracy: 0.96 ± 0.02

```
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
skf_scores = []
for train_index, test_index in skf.split(X, y):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)
    score = accuracy_score(y_test, predictions)
    skf_scores.append(score)
print(f'Stratified K-Fold Accuracy: {np.mean(skf_scores):.2f} ± {np.std(skf_scores):.2f}')
```

Stratified K-Fold Accuracy: 0.95 ± 0.03

7.2 Systematically explore combinations of hyperparameters to optimize model performance.(use grid and randomized search)

```
import numpy as np
import pandas as pd

from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split, GridSearchCV, Random
izedSearchCV
from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score

data = load_iris()
X, y = data.data, data.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
random_state=42)
model = RandomForestClassifier(random_state=42)
param_grid = {
    'n_estimators': [10, 50, 100],
    'max_depth': [None, 5, 10, 20],
    'min_samples_split': [2, 5, 10],
}

grid_search = GridSearchCV(estimator=model, param_grid=param_grid, scoring
='accuracy', cv=5)
grid_search.fit(X_train, y_train)
print("Grid Search Best Parameters:", grid_search.best_params_)
print("Grid Search Best Score:", grid_search.best_score_)

Grid Search Best Parameters: {'max_depth': 5, 'min_samples_split': 5, 'n_e
stimators': 10}
Grid Search Best Score: 0.9636363636363636

param_dist = {
    'n_estimators': np.arange(10, 200, 10),
    'max_depth': [None] + list(np.arange(1, 20, 1)),
    'min_samples_split': np.arange(2, 20, 2),
}

random_search = RandomizedSearchCV(estimator=model, param_distributions=pa
ram_dist, n_iter=50, scoring='accuracy', cv=5, random_state=42)
random_search.fit(X_train, y_train)
print("Randomized Search Best Parameters:", random_search.best_params_)
print("Randomized Search Best Score:", random_search.best_score_)

Randomized Search Best Parameters: {'n_estimators': 120, 'min_samples_spli
```

```
t': 16, 'max_depth': None}  
Randomized Search Best Score: 0.9636363636363636
```

```
best_model = grid_search.best_estimator_  
y_pred = best_model.predict(X_test)
```

Practical 8

8.1 Implement Bayesian Learning using inferences

```
import numpy as np

P_A = 0.5
P_B = 0.5

def likelihood_heads(coin, flips):
    if coin == 'A':
        return (0.5 * flips) * (0.5 * (10 - flips))
    elif coin == 'B':
        return (0.9 * flips) * (0.1 * (10 - flips))

observed_heads = 8
total_flips = 10

likelihood_A = likelihood_heads('A', observed_heads)
likelihood_B = likelihood_heads('B', observed_heads)

marginal_likelihood = (likelihood_A * P_A) + (likelihood_B * P_B)

posterior_A = (likelihood_A * P_A) / marginal_likelihood
posterior_B = (likelihood_B * P_B) / marginal_likelihood

print(f"Posterior Probability of Coin A: {posterior_A:.4f}")
print(f"Posterior Probability of Coin B: {posterior_B:.4f}")

Posterior Probability of Coin A: 0.7353
Posterior Probability of Coin B: 0.2647
```

Practical 9

9.1 Set up a generator network to produce samples and a discriminator network to distinguish between real and generated data. (Use a simple small dataset)

```
import numpy as np

import matplotlib.pyplot as plt

import tensorflow as tf

from tensorflow.keras import layers, models

(X_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
X_train = (X_train.astype(np.float32) - 127.5) / 127.5
X_train = np.expand_dims(X_train, axis=-1)

latent_dim = 100
num_examples_to_generate = 16

def build_generator():
    model = models.Sequential()
    model.add(layers.Dense(256, input_dim=latent_dim))
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.BatchNormalization())
    model.add(layers.Dense(512))
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.BatchNormalization())
    model.add(layers.Dense(1024))
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.BatchNormalization())
    model.add(layers.Dense(28 * 28 * 1, activation='tanh'))
    model.add(layers.Reshape((28, 28, 1)))
    return model

def build_discriminator():
    model = models.Sequential()
    model.add(layers.Flatten(input_shape=(28, 28, 1)))
    model.add(layers.Dense(512))
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.Dense(256))
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.Dense(1, activation='sigmoid'))
    return model
```

```
generator = build_generator()  
discriminator = build_discriminator()  
discriminator.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```



```

discriminator.trainable = False
gan_input = layers.Input(shape=(latent_dim,))
generated_image = generator(gan_input)
gan_output = discriminator(generated_image)
gan = models.Model(gan_input, gan_output)
gan.compile(optimizer='adam', loss='binary_crossentropy')

def generate_and_save_images(model, epoch, test_input):
    predictions = model(test_input)
    predictions = (predictions.numpy() + 1) / 2 # Rescale to [0, 1]

    plt.figure(figsize=(4, 4))
    for i in range(predictions.shape[0]):
        plt.subplot(4, 4, i + 1)
        plt.imshow(predictions[i, :, :, 0], cmap='gray')
        plt.axis('off')
    plt.savefig(f'gan_epoch_{epoch}.png')
    plt.show()

def train_gan(epochs, batch_size):
    random_latent_vectors = tf.random.normal(shape=(num_examples_to_generate, latent_dim))

    for epoch in range(epochs):

        idx = np.random.randint(0, X_train.shape[0], batch_size)
        real_images = X_train[idx]

        noise = tf.random.normal(shape=(batch_size, latent_dim))
        fake_images = generator(noise)

        combined_images = tf.concat([real_images, fake_images], axis=0)

        labels = tf.constant([[1.0]] * batch_size + [[0.0]] * batch_size)

        d_loss = discriminator.train_on_batch(combined_images, labels)

        noise = tf.random.normal(shape=(batch_size, latent_dim))
        misleading_labels = tf.constant([[1.0]] * batch_size)

        g_loss = gan.train_on_batch(noise, misleading_labels)

        if epoch % 100 == 0:
            print(f"Epoch: {epoch}")
            print(f"Discriminator Loss: {d_loss[0]}")
            print(f"Generator Loss: {g_loss}")

        generate_and_save_images(generator, epoch, random_latent_vectors)
rs)

```


```
epochs = 300  
batch_size = 64
```

```
train_gan(epochs, batch_size)
```

Epoch: 0

Discriminator Loss: 0.7258248329162598


Generator Loss: [array(0.72582483, dtype=float32), array(0.72582483, dtype=float32), array(0.390625, dtype=float32)]



Epoch: 100

Discriminator Loss: 2.1150412559509277


Generator Loss: [array(2.1150413, dtype=float32), array(2.1150413, dtype=float32), array(0.20482673, dtype=float32)]



Epoch: 200

Discriminator Loss: 2.8626105785369873

Generator Loss: [array(2.8626106, dtype=float32), array(2.8626106, dtype=float32), array(0.20747824, dtype=float32)]



Practical 10

10.1 Develop an API to deploy your model and perform predictions

Required Libraries

```
!pip install pyngrok flask scikit-learn
```

Importing Libraries

```
from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
import pickle

from flask import Flask, request, jsonify

from pyngrok import ngrok
```

Load dataset

```
iris = load_iris()
X, y = iris.data, iris.target
```

Split data

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Train a model

```
model = RandomForestClassifier()
model.fit(X_train, y_train)
```

Save the model

```
with open('model.pkl', 'wb') as model_file:
    pickle.dump(model, model_file)
```

Load the model

```
with open('model.pkl', 'rb') as model_file:
    model = pickle.load(model_file)
```

Create Flask app app =

```
Flask(__name__)
port = "5000"
```

```
@app.route('/')
def home():

    return "Welcome to the Iris Prediction API! Use the /predict endpoint
to make predictions."

@app.route('/predict', methods=['POST'])
def predict():

    data = request.json
    features = data.get('features')
```

```

# Ensure the features are in the correct format

if not features or len(features) != 4: # Assuming 4 features for iris dataset

    return jsonify({'error': 'Invalid input format. Please provide 4 features.'}), 400

try:

    prediction = model.predict([features]) # Wrap features in a list
to create 2D array

    return jsonify({'prediction': int(prediction[0])}) # Convert prediction to int

except Exception as e:

    return jsonify({'error': str(e)}), 500

# Start ngrok and print the public URL
ngrok.set_auth_token("api_auth_token") public_url =
ngrok.connect(port).public_url print("Public URL:", public_url)

# Run the Flask app

if __name__ == '__main__':
    app.run(port=port)

```

```

Requirement already satisfied: pyngrok in /usr/local/lib/python3.10/dist-packages (7.2.0)
Requirement already satisfied: flask in /usr/local/lib/python3.10/dist-packages (2.2.5)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.5.2)
Requirement already satisfied: PyYAML>=5.1 in /usr/local/lib/python3.10/dist-packages (from pyngrok) (6.0.2)
Requirement already satisfied: Werkzeug>=2.2.2 in /usr/local/lib/python3.10/dist-packages (from flask) (3.0.4)
Requirement already satisfied: Jinja2>=3.0 in /usr/local/lib/python3.10/dist-packages (from flask) (3.1.4)
Requirement already satisfied: itsdangerous>=2.0 in /usr/local/lib/python3.10/dist-packages (from flask) (2.2.0)
Requirement already satisfied: click>=8.0 in /usr/local/lib/python3.10/dist-packages (from flask) (8.1.7)
Requirement already satisfied: numpy>=1.19.5 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.26.4)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.13.1)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.5.0)

```

Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2>=3.0->flask) (2.1.5)
Public URL: <https://2f62-49-43-24-101.ngrok-free.app>

* Serving Flask app '__main__'
* Debug mode: off

INFO:werkzeug:WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.

* Running on http://127.0.0.1:5000

INFO:werkzeug:Press CTRL+C to quit

INFO:werkzeug:127.0.0.1 - - [24/Oct/2024 07:50:42] "POST /predict HTTP/1.1"
200 -

The screenshot shows a web browser window with a POST request to the endpoint `https://2f62-49-43-24-101.ngrok-free.app/predict`. The request body is a JSON object with a `features` array containing the values `5.1`, `3.5`, `1.4`, and `0.2`. The response status is `200 OK` with a response time of `673 ms` and a response size of `195 B`. The response body is a JSON object with a `prediction` value of `0`.

POST `https://2f62-49-43-24-101.ngrok-free.app/predict`

Params Authorization Headers (9) Body Pre-request Script Tests Settings Cookies Beautify

none form-data x-www-form-urlencoded raw binary JSON

```
1  {
2    "features": [
3      5.1,
4      3.5,
5      1.4,
6      0.2
7    ]
8  }
```

Body Cookies Headers (5) Test Results 200 OK 673 ms 195 B Save Response

Pretty Raw Preview Visualize JSON

```
1  {
2    "prediction": 0
3  }
```